Refinement versus Refresh: Compute Allocation for Closed-Loop Robot Policies

Alba María Téllez Fernández
MAAT Institute
alba.tellez@maatinstitute.com

Abstract

As vision–language–action policies become more computationally demanding, deployment under explicit resource constraints poses an allocation problem: how much inference effort should construct each action chunk, and how long should it execute before prediction is refreshed from a new observation? We study this question by varying generation effort and execution prefix separately, measuring both closed-loop success and complete inference expenditure, and distinguishing prespecified comparisons from later descriptive analyses. We use one frozen SmolVLA checkpoint in synchronous LIBERO simulation as a worked case, with denoising count as its generation control. At a fixed two-action prefix, the prespecified comparison found 139/160 successes with two denoising steps versus 122/160 with five, while using 5.9% less reference-calibrated arithmetic per control step; 82.1% of native generation arithmetic was independent of denoising. A later descriptive contrast found 138/160 versus 139/160 successes when executing five rather than two actions per prediction, with 60.7% lower arithmetic per attempt; equivalence was not established. These checkpoint-specific results do not define a general allocation rule. They show why closed-loop behavior, complete inference cost, and the status of each comparison must be considered together.

1 Introduction

Deploying vision–language–action policies under a finite inference budget is an allocation problem. A chunked policy spends internal computation constructing an action sequence, executes some prefix of that sequence, and then pays for another prediction conditioned on a newer observation. More generation effort may change the predicted chunk; a longer execution prefix reduces prediction frequency but delays feedback and exposes later action positions. Comparing these controls therefore requires closed-loop outcomes and a common resource boundary, not generation cost alone.

We ask: *given a fixed inference budget, how should a chunked policy allocate compute between action generation and observation-conditioned execution?* To answer this without turning one architecture's control into a general rule, we compare the model-specific generation-effort setting and execution prefix separately. We measure complete prediction and realized episode expenditure, separate discovery from held-out evaluation, freeze confirmatory comparisons before outcomes are observed, and mark later analyses as descriptive. This combination of controlled comparison, complete cost accounting, and explicit evidence status is the primary transferable contribution.

The comparison is intended for chunked policies that expose a controllable generation effort. That control could be the number of integration steps in an iterative flow or diffusion policy, or the candidate, search, or verification effort in a policy that uses such mechanisms. Only denoising count is tested here, and the controls need not have comparable units across architectures. We use

40th Conference on Neural Information Processing Systems (NeurIPS 2026). Workshop: Post-Training Adaptation of Robot Foundation Models.

synchronous execution to hold sequencing and measurement boundaries fixed; the study does not identify an asynchronous scheduling optimum, real-robot latency, or a universal hardware winner.

SmolVLA provides the worked case study [1]. Its generation control is denoising count, and its execution control is the number of actions applied before replanning. We evaluate twelve configurations in a paired 480-attempt discovery study, then seven configurations in 1,120 held-out attempts on eight tasks excluded from selection. Three denoising comparisons were frozen before confirmation. The execution-prefix contrasts were added later and remain descriptive; they neither isolate observation freshness from action-position effects nor establish an optimal allocation rule.

For this checkpoint, the prespecified primary comparison yields 139/160 successes with two denoising steps versus 122/160 with five at the same two-action prefix, while using less counted arithmetic. Calibration shows that 82.1% of native generation arithmetic is independent of denoising, so reducing refinement alone leaves substantial fixed work. At two denoising steps, executing five actions per prediction yields 138/160 successes and 60.7% less arithmetic per attempt than executing two; equivalent quality is not established. A simple cost-aware rule also matches the fitted response model's choices at all three ceilings, with no demonstrated selection advantage from fitting.

These checkpoint-specific outcomes show both what the comparison can reveal and where its evidence limits the resulting claims.

The contributions are:

- a formulation that relates generation effort and execution prefix to complete prediction and episode expenditure;
- a controlled evaluation that distinguishes prespecified tests from amended descriptive contrasts and retains their uncertainty and provenance; and
- a paired SmolVLA case study, using denoising count as the generation control, across twelve discovery configurations and seven held-out settings.

2 Related Work

Resource-allocation studies motivate evaluating configurations across explicit budgets. Chinchilla studies model-size and training-data allocation [2]; embodied scaling studies consider world modeling, behavior cloning, and closed-loop forecasting and planning [3, 4]. Our comparison instead varies deployment refinement and execution prefix within a fixed policy. Work distinguishing FLOP-optimal from elapsed- or monetary-cost-optimal allocation [5] also motivates retaining separate resource coordinates.

VLA-Perf characterizes inference performance across denoising, chunk size, asynchronous execution, and hardware/runtime choices [6]. SmolVLA studies asynchronous inference and robot behavior [1]. Here predicted chunk size remains fixed, while executed prefix changes; task success is measured alongside computation. RoboMonkey studies test-time sampling and action verification [7], an adjacent allocation axis that would require counting candidate generation and verification costs.

Adaptive VLA scheduling is directly related. DVAC uses denoising-derived action variance to choose an execution prefix before replanning [8]. ElegantVLA schedules vision–language recomputation or reuse together with action denoising modes [9]. Our study instead compares fixed configurations under explicit arithmetic ceilings and freezes separate-task denoising comparisons. It neither introduces an adaptive scheduler nor establishes an advantage over these methods; their reported speedups have different policies, hardware, and measurement boundaries.

3 Allocating Compute Between Generation and Execution

The evaluation has four stages. First, choose an architecture-specific generation effort control r and an execution prefix h, while holding the checkpoint and paired rollout conditions fixed. Here r denotes a controllable amount of internal inference work used to produce one chunk; its implementation depends on the policy. Second, measure complete generation cost and realized episode expenditure under an explicit counting boundary. Third, use discovery data to freeze confirmatory comparisons before held-out outcomes are observed. Fourth, report success and expenditure together, label later contrasts as descriptive, and avoid interpreting arithmetic savings as latency improvement or similar

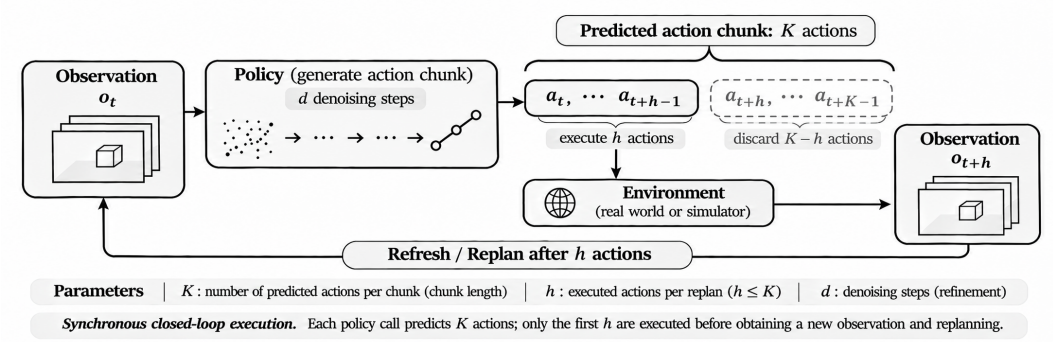


Figure 1: Closed-loop execution in the SmolVLA case study. Each policy call predicts K actions after d denoising steps, executes the first h, discards the remainder, and replans from a new observation. Here $r = d$ and $K = 50$.

observed success as equivalent quality. The following subsections instantiate each stage with $r = d$, the denoising count of SmolVLA.

3.1 Compute rate and episode expenditure

Let K be the predicted action-chunk length, $h < K$ the execution prefix, and r the architecture-specific generation-effort control. At each policy query, the policy predicts K actions, executes only the first h, discards the remaining $K - h$, and generates a new chunk from a new observation. Thus h controls replanning frequency without changing K. We reserve T_e for the number of actions executed in attempt e, and write $F_{\text{gen}}(r)$ for the complete measured arithmetic of one generation under a stated counting boundary. The study compares configurations (h, r) within a policy; it does not assume that numerical values of r are comparable across architectures. For complete execution prefixes, the nominal compute rate per simulated control step is

$$b(h, r) = \frac{F_{\text{gen}}(r)}{h}.$$

For a candidate set $\mathcal{C}$ and budget ceiling B, the allocation target is

$$(h^\star, r^\star)(B) \in \arg \max_{(h,r)\in\mathcal{C}:\, b(h,r)\le B} p_{\text{success}}(h, r). \tag{1}$$

A configuration need not exhaust the ceiling. This is a deployment compute-rate constraint for a fixed checkpoint, rather than a training-plus-deployment objective.

In the SmolVLA instantiation, $K = 50$ and $r = d$, where d is the denoising count. Figure 1 shows this case-study execution loop. Each policy call predicts a fixed-length chunk, executes only the first h actions, discards the remainder, and refreshes from the resulting observation. The denoising count controls generation effort inside the policy call, while h controls how frequently the policy receives a new observation.

We distinguish this rate from total arithmetic per attempted episode. If attempt e executes T_e actions, its generated-chunk expenditure is

$$C_e = \lceil T_e/h \rceil \, F_{\text{gen}}(r). \tag{2}$$

The final generated chunk is charged even if only part of its execution prefix is used. All-attempt means include unsuccessful episodes. Different trajectories and termination times can therefore change expenditure at the same nominal rate. FLOPs, host elapsed time, and CUDA-stream time retain separate units and measurement boundaries.

3.2 The fixed cost of a fresh prediction

Complete-generation profiling counts matrix products, convolutions, and attention across vision, language prefill, and the action expert, using two FLOPs per multiply-add. Embedding lookups, normalization, activations, elementwise arithmetic, and external preprocessing are excluded. We

therefore refer to *dense neural arithmetic*, not an exact count of every scalar operation. Calibration uses one real development observation and fixed tensor shapes; its calibration procedure is detailed in Appendix A.

For the case-study checkpoint and reference input, the measured arithmetic satisfies, to rounding,

$$F_{\text{gen}}(d) \simeq 527.108 + 11.457d \quad \text{GFLOPs}$$

at all six calibrated denoising settings. Native generation at $d = 10$ costs 641.675 GFLOPs, of which 82.1% is independent of d. Halving denoising from ten to five saves 8.9% of counted generation work; reducing it to two saves 14.3%. Equivalently, relative to native deployment $(h, d) = (1, 10)$,

$$\frac{b(h,d)}{b_{\text{native}}} \simeq \frac{0.821 + 0.179(d/10)}{h}.$$

We prespecify compute-rate ceilings of 0.25, 0.50, and 1.00 times native. Eligibility uses measured per-d values, not the rounded affine expression. For example, $(1, 2)$ costs 0.857 times native rate and is ineligible at half-native budget, whereas $(2, 2)$ costs 0.429 times native. The denoising-independent work alone excludes $h = 1$ at half-native budget and $h \leq 2$ at quarter-native budget. A tighter budget can therefore force a longer execution prefix before any quality-based choice is made. This is an implementation-specific arithmetic floor under the declared counting boundary, not a universal VLA scaling law.

Calibration uses a fixed input shape; a separate language-length audit leaves eligibility unchanged under the original absolute ceilings (Appendix A.3). Reference arithmetic remains distinct from task-shaped arithmetic and elapsed time. For the case-study implementation on a fixed input, reducing d from 10 to 5 lowers measured host generation time by 40.5% while counted arithmetic falls by 8.9% (Appendix A.1); FLOPs are therefore not a latency proxy. The cost model defines feasible schedules; the following experiments ask which feasible expenditures help closed-loop control.

4 SmolVLA Case Study

4.1 Discovery and confirmation design

We evaluate the public `HuggingFaceVLA/smolvla_libero` checkpoint on LIBERO [10]. The predicted chunk length remains fixed at $K = 50$; native execution uses $(h, d) = (1, 10)$. We hold weights, sensors, preprocessing, action normalization, control period, and task-specific episode horizons fixed.

Discovery crosses $h \in \{1, 2, 5\}$ and $d \in \{2, 5, 10, 20\}$ on Spatial tasks 5–8, five initial states per task, and two policy seeds. Each of the 40 paired task–state–seed identities receives all twelve configurations, giving 480 attempts. Tasks 5 and 6 overlap the pilot; tasks 7 and 8 were selected before observing their outcomes in this study. Condition order is randomized within each paired identity.

Discovery is used to freeze three budget-specific denoising comparisons before confirmation. Their selected settings and comparators, together with native, give seven configurations. Confirmation evaluates all seven on Goal and Object tasks 5–8, ten states per task, and two seeds: 160 attempts per configuration, totaling 1,120. These eight tasks are excluded from configuration selection; prior exposure during unrelated development is not assessed. The sample allocation prioritizes task and state coverage; no quality-preservation margin is prespecified. Appendix A.5 distinguishes the original tests from the amended descriptive comparisons.

Initial simulator states and observations match across conditions. Policy noise is keyed by task, state, seed, and simulation tick, preserving a common noise key whenever conditions generate at the same tick. Solver timesteps and a constant-field integration check verify that every d traverses the full interval from 1 to 0; fewer steps change the discretization, not the integration endpoint. Recorded initial conditions, generated and executed actions, and outcomes agree across the evaluation records. LIBERO's sparse reward and success indicator derive from the same task-goal evaluation, so their agreement checks record consistency rather than supplying an independent outcome measure. The recorded outcome path is traced to the benchmark's task-specific goal conditions. These checks establish strong internal consistency rather than an independent simulator replay (Appendix A.5).

Success averages both seeds within each state, then states and tasks with equal task weights. Descriptive 95% intervals use 20,000 bootstrap draws of states within the fixed tasks, retaining paired seeds and configurations. They condition on four discovery or eight confirmation tasks; secondary intervals are unadjusted. Arithmetic includes every generated chunk, including failures and partly consumed final prefixes. Timing measurements and their monitoring limits are reported separately in Appendix A.1.

4.2 Discovery: quality does not increase monotonically with refinement

In discovery for the case-study checkpoint, native $(1, 10)$ succeeds in 25/40 attempts. The highest observed count is 28/40 for $(1, 5)$; $(1, 2)$ and $(2, 2)$ each succeed in 26/40. At $h = 5$, successes are 21, 18, 16, and 14 out of 40 as d increases from 2 to 20. Appendix Figure 3 shows every configuration. These point estimates do not justify imposing monotonic improvement with denoising. The small paired sample also does not certify equal quality for configurations with similar observed counts.

The resource coordinate also matters: nominal rate differs from realized episode expenditure, and arithmetic savings differ from elapsed savings. Substituting a host-time constraint changes feasible refresh choices; the appendix reports this sensitivity analysis separately from the discovery success figure. It does not change the frozen arithmetic-based experiment. Full plots, task counts, and resource tables appear in Appendix A.

4.3 Frozen tests and the role of selection

For this case study, a response model fitted only to discovery selects $(5, 2)$, $(2, 2)$, and $(1, 2)$ at the quarter-, half-, and full-native ceilings. Their frozen comparators are $(5, 5)$, $(2, 5)$, and $(1, 5)$, respectively; the half-native comparison is primary. All selected settings are existing discovery grid points, so confirmation tests task transfer rather than extrapolation to an untested configuration. Each pair changes denoising at a fixed prefix. These are comparisons under common ceilings, not equal-rate tests of how to divide compute between refinement and refresh.

Within this case study, $(2, 2)$ has the highest eligible discovery success at the primary ceiling. Its comparator $(2, 5)$ ties with $(2, 10)$ for the highest count among distinct eligible alternatives and wins the lower-cost tie-break. A simple rule choosing minimum denoising and then the shortest feasible prefix makes the same choices as the fitted model at all three ceilings. This baseline was documented after collection began and is interpreted descriptively. Transfer outcomes cannot demonstrate a selection advantage between rules that make identical choices. Model fitting, the full candidate set, and selector diagnostics are retained in Appendix B.

5 SmolVLA Worked Example: Held-Out Results

5.1 Complete confirmation overview

All 1,120 originally planned SmolVLA case-study confirmation attempts are complete. Table 1 reports the seven configurations on the same paired identities. The audit excludes two repeated completed records under the documented precedence rule; continuation completes the original design, not an independent replication (Appendix A). We first report the frozen denoising tests, then the amended prefix comparisons. Only the former support confirmatory claims; the latter are reported as descriptive analyses with their provenance and uncertainty retained.

Table 1: SmolVLA case-study confirmation: 160 attempts per configuration on eight tasks. $K = 50$ is the fixed predicted chunk length, h is the executed prefix, and d is the denoising count. Nominal rate/native is $F_{\text{gen}}(d)/(hF_{\text{gen}}(10))$. All-attempt arithmetic uses the length-20 reference calibration and charges every generation. The primary pair is $(2, 2)$ versus $(2, 5)$.

Allocation		Success		Nominal rate	Mean per attempt	
h	d	Count	(%)	/ native	TFLOPs	Steps
1	2	137/160	85.6	0.857	83.3	151.4
1	5	136/160	85.0	0.911	89.1	152.5
1	10	135/160	84.4	1.000	98.4	153.4
2	2	139/160	86.9	0.429	41.1	149.0
2	5	122/160	76.3	0.455	48.7	166.3
5	2	138/160	86.3	0.171	16.1	144.9
5	5	133/160	83.1	0.182	18.5	156.2

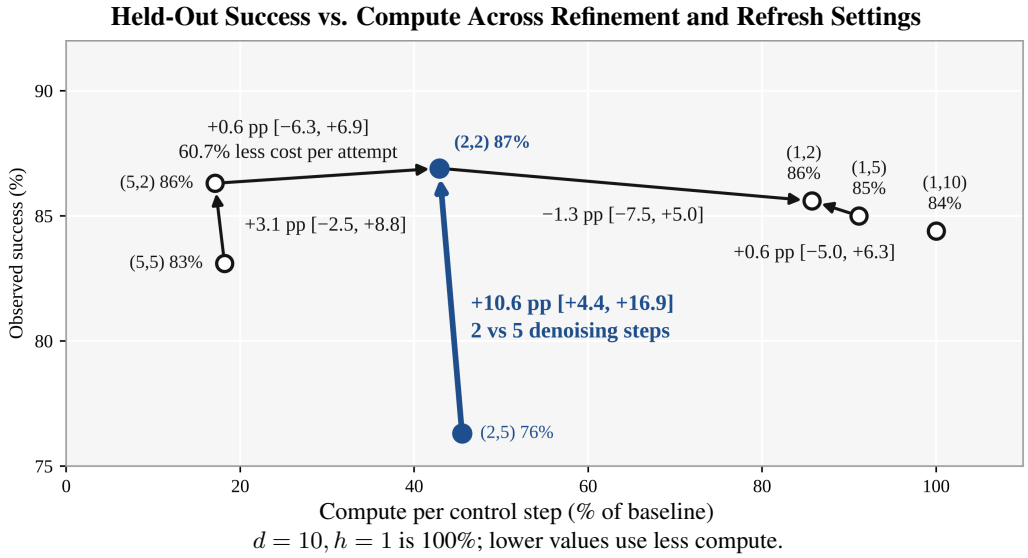


Figure 2: SmolVLA case-study held-out success, compute, and paired comparisons. Each point is one of the seven confirmation settings, labelled (h, d) for executed prefix and denoising count, and aggregates 160 attempts over eight tasks. The vertical axis gives observed success and begins at 75%. The horizontal axis gives nominal counted generation arithmetic per control step as a percentage of the native $(d = 10, h = 1)$ deployment baseline; leftward means less counted work per executed action. Arrows run from comparator B to setting A and print the A-minus-B success difference with its task-conditional, descriptive 95% paired state-bootstrap interval. Blue identifies the prespecified $(2, 2)$ versus $(2, 5)$ refinement test. Black arrows show two other frozen refinement comparisons and two later-amended descriptive contrasts. The 60.7% annotation is the separate reduction in mean arithmetic per attempted episode for $(5, 2)$ relative to $(2, 2)$; the associated $+0.6$-point success contrast does not establish equivalent quality. The two additional amended contrasts involving $(1, 10)$, complete task-level outcomes, and full episode expenditures are reported in Appendix A.5 and Table 1.

Table 2: Evidence status and permitted interpretation. Effects use A minus B. The table separates prespecified confirmation from amended descriptive analysis.

Comparison	Status	Result	Supported claim
$(2,2)-(2,5)$	Frozen primary	$+10.6$ pp; 95% interval $[+4.4,+16.9]$	At $h=2$, fewer denoising steps improved observed success for this checkpoint.
Other fixed-h refinement tests	Frozen	Positive point estimates; intervals span zero	A uniform benefit is not established for this checkpoint.
$(2,2)-(5,2)$	Amended descriptive	$+0.6$ pp; 95% interval $[-6.3,+6.9]$; $(5,2)$ uses 60.7% less arithmetic per attempt	The case study shows large cost savings; equivalent quality is unresolved.

5.2 Refinement: the frozen denoising comparisons

For the case-study checkpoint, the primary two-action prefix with two denoising steps yields 139/160 successes versus 122/160 for five steps. The paired difference is $+10.6$ percentage points, with descriptive 95% interval $[+4.4,+16.9]$, conditional on the eight tasks. The selected setting also uses 5.9% less nominal reference arithmetic per control step: 275.0 versus 292.2 GFLOPs. Thus, in this SmolVLA comparison, reducing refinement improves observed success while reducing counted work at this fixed prefix.

Mean reference arithmetic per attempt falls from 48.7 to 41.1 TFLOPs (15.6%). This larger reduction includes changed trajectories and termination: mean episode length is 149.0 versus 166.3 control steps. All failures reach the task horizon (21 versus 38); the remaining attempts terminate successfully. The primary direction is positive on seven tasks and tied on one, with 24 paired wins and seven losses. Appendix A.4 gives task sensitivity and expenditure intervals. These measurements do not establish a controlled physical-latency improvement. Figure 2 places the primary setting in the joint success–compute coordinate and places its interval alongside all seven confirmation contrasts.

Within the same checkpoint, the other two frozen tests have positive point estimates but intervals spanning zero (Figure 2). At $h=5$, two versus five denoising steps give 138/160 versus 133/160 successes; at $h=1$, they give 137/160 versus 136/160. The largest observed benefit is at the primary prefix. The primary gap exceeds its frozen prediction, while matching the signs of three point estimates does not establish calibrated effect-size prediction. These results do not demonstrate a uniform denoising benefit across prefixes.

5.3 Refresh: descriptive prefix comparisons

Within the SmolVLA case study, we now hold denoising fixed and compare expenditures on different refresh schedules. Table 2 records the later amendment and the claim boundary; Appendix A.5 gives full provenance.

At $d=2$, $(2,2)$ exceeds $(5,2)$ by just one success: 139/160 versus 138/160, a paired difference of $+0.6$ percentage points $[-6.3,+6.9]$. Yet $(5,2)$ uses 60% less nominal arithmetic per control step and 60.7% less reference arithmetic per attempt: 16.1 versus 41.1 TFLOPs. The interval permits meaningful differences in either direction. At task level, the $(5,2)-(2,2)$ difference is positive on three tasks, negative on three, and tied on two (Table 9). Thus the near-equal aggregate counts contain mixed descriptive task-level changes.

Refreshing every action with $(1,2)$ instead of every two actions gives two fewer successes, a difference of -1.3 points $[-7.5,+5.0]$, at 83.3 versus 41.1 TFLOPs per attempt. More frequent refresh has no demonstrated quality benefit in these SmolVLA contrasts. Native-relative comparisons, which also change denoising, are reported separately in Appendix A.5.

Within the case study, the discovery point estimates favor shorter prefixes at $d=2$ (Section 4.2), but their quality advantage is not established on transfer. The task populations differ, so this comparison does not identify a causal task effect. The cost model explains feasibility, but does not determine the quality ranking on new tasks.

6 Discussion and Conclusion

This study asks how a chunked robot policy should divide inference compute between constructing an action chunk and deciding how long to execute it before generation is refreshed from a new observation. Answering that question requires complete policy-call cost, held-out closed-loop outcomes, and a clear distinction between prespecified and later comparisons. Together, these prevent a cheap setting from being treated as useful without behavioral evidence, or a successful setting from being discussed without its compute cost.

The broader methodological proposal is that this comparison can be defined across policy classes by treating generation effort as an architecture-specific control. For diffusion or flow policies it may be an integration-step count; for sampling, search, or verification systems it may be the number of candidates, search expansions, or verification passes, provided the control and its complete cost are defined in advance. SmolVLA serves as a worked case study for one checkpoint and simulator stack, demonstrating the comparison with denoising count and showing how evidence status changes the permitted claims. In this SmolVLA case study, additional refinement did not consistently produce better closed-loop behavior. Reducing refinement alone delivered limited savings because most generation work was independent of denoising. Changing the execution prefix affected cost more strongly, but similar observed success was insufficient to establish equivalent quality. In the case study, the fitted response model also did not improve configuration choices over a simple cost-aware rule. Generation effort and execution prefix should therefore be evaluated together within each model, under explicit counting boundaries, with frozen confirmatory tests visibly separated from amended descriptive analyses. These outcomes remain checkpoint-specific; broader allocation claims require the same within-model study across other checkpoints, policy architectures, tasks, and execution stacks.

Interpretation and limits. The two deployment controls are not isolated causal interventions on prediction quality and observation freshness. Changing denoising changes numerical integration; changing prefix changes both observation age and executed action positions. The mechanism behind the observed primary difference remains open. The study covers one checkpoint and twelve fixed tasks in simulation, with only four discovery task-level validation folds. Intervals are task-conditional, secondary contrasts are unadjusted, and no quality-preservation margin was prespecified. The refresh-prefix contrasts were added after individual outcomes were visible and remain descriptive. Two steps is the minimum denoising tested; $d = 1$ is untested. The arithmetic boundary excludes several operation classes; language lengths are audited but image shape is fixed. Sampled GPU checks have coverage gaps, and simulated control steps do not establish real-time responsiveness. The recorded outcomes trace to LIBERO's task-goal evaluation, and one native-setting parity episode reproduces the reference evaluator's initial state, action sequence, episode length, and outcome. However, the confirmation actions were not replayed in a separate, version-matched simulator. Outcome validity is therefore strongly supported internally but not fully independently established.

Next steps. A controlled stale-input diagnostic is needed to separate observation freshness from the action-position effects introduced by changing the executed prefix. A predeclared quality-preservation test should evaluate the cheaper $(5, 2)$ setting, and independent outcome replay should test the retained primary trajectories in a separate, version-matched simulator. The same study design can then be applied within other policies and execution stacks: define a controllable generation-effort axis and execution prefix, calibrate complete inference cost, freeze within-policy comparisons, and evaluate success and expenditure jointly. Cross-architecture evidence should compare the within-model allocation results rather than pool unlike control values or counting boundaries. Replication across checkpoints, benchmark families, and asynchronous or physical-robot execution is required to determine when additional generation work is worth paying for and when compute is better spent producing a new chunk from fresher feedback.

References

[1] Mustafa Shukor, Dana Aubakirova, Francesco Capuano, Pepijn Kooijmans, Steven Palma, Adil Zouitine, Michel Aractingi, Caroline Pascal, Martino Russi, Andres Marafioti, Simon Alibert, Matthieu Cord, Thomas Wolf, and Remi Cadene. SmolVLA: A vision-language-action model for affordable and efficient robotics. arXiv preprint arXiv:2506.01844v1, 2025. URL https://arxiv.org/abs/2506.01844v1.

[2] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Thomas Hennigan, Eric Noland, Katherine Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karén Simonyan, Erich Elsen, Oriol Vinyals, Jack Rae, and Laurent Sifre. An empirical analysis of compute-optimal large language model training. In *Advances in Neural Information Processing Systems*, volume 35, 2022. doi: 10.52202/068431-2176. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/c1e2faff6f588870935f114ebe04a3e5-Abstract.html.

[3] Tim Pearce, Tabish Rashid, David Bignell, Raluca Georgescu, Sam Devlin, and Katja Hofmann. Scaling laws for pre-training agents and world models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48542–48562. PMLR, 2025. URL https://proceedings.mlr.press/v267/pearce25a.html.

[4] Mustafa Baniodeh, Kratarth Goel, Scott Ettinger, Carlos Fuertes, Ari Seff, Tim Shen, Cole Gulino, Chenjie Yang, Ghassen Jerfel, Dokook Choe, Rui Wang, Benjamin Charrow, Vinutha Kallem, Sergio Casas, Rami Al-Rfou, Benjamin Sapp, and Dragomir Anguelov. Scaling laws of motion forecasting and planning – technical report. arXiv preprint arXiv:2506.08228v2, 2025. URL https://arxiv.org/abs/2506.08228v2.

[5] Nikhil Sardana, Jacob Portes, Sasha Doubov, and Jonathan Frankle. Beyond chinchilla-optimal: Accounting for inference in language model scaling laws. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 43445–43460. PMLR, 2024. URL https://proceedings.mlr.press/v235/sardana24a.html.

[6] Wenqi Jiang, Jason Clemons, Karu Sankaralingam, and Christos Kozyrakis. How fast can I run my VLA? demystifying VLA inference performance with VLA-Perf. arXiv preprint arXiv:2602.18397v1, 2026. URL https://arxiv.org/abs/2602.18397v1.

[7] Jacky Kwok, Christopher Agia, Rohan Sinha, Matt Foutter, Shulu Li, Ion Stoica, Azalia Mirhoseini, and Marco Pavone. RoboMonkey: Scaling test-time sampling and verification for vision-language-action models. In *Proceedings of the 9th Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 3200–3217. PMLR, 2025. URL https://proceedings.mlr.press/v305/kwok25a.html.

[8] Xiangdong Feng, Yuxuan Cheng, Chen Shi, Boyao Han, Yuxuan Yan, Yitong Hong, Zhuotao Tian, and Li Jiang. Denoising tells when to replan: Denoising-variance adaptive chunking for flow-based robot policies. arXiv preprint arXiv:2606.03847v1, 2026. URL https://arxiv.org/abs/2606.03847v1.

[9] Ye Li, Huanan Liu, Kangye Ji, Yuan Meng, Jiajun Fan, Yuansong Wang, Shiyu Qin, Chenglei Wu, Shu-Tao Xia, and Zhi Wang. ElegantVLA: Learning when to think for efficient vision-language-action models. arXiv preprint arXiv:2605.29438v1, 2026. URL https://arxiv.org/abs/2605.29438v1.

[10] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. LIBERO: Benchmarking knowledge transfer for lifelong robot learning. In *Advances in Neural Information Processing Systems*, volume 36, 2023. doi: 10.52202/075280-1939. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets_and_Benchmarks.html.

A Detailed Outcomes and Resource Measurements

The checkpoint is `HuggingFaceVLA/smolvla_libero`, revision `6721902bc4d61e50a3bfdb11dfb4cb626f05d102`. Predicted chunk size is 50; native execution uses $h = 1, d = 10$. Discovery uses Spatial tasks 5–8, initial states 10–14, and policy seeds 2000 and 2001. The 20 task–state blocks contain both seeds and all 12 configurations. Confirmation uses Goal and Object tasks 5–8, states 20–29, and policy seeds 3000 and 3001. Its 80 task–state blocks each contain both seeds and all seven configurations. Discovery bootstrap resampling uses 20,000 draws with analysis seed 20260913.

Table 3: Discovery successes on Spatial tasks 5–8. Each task contributes five states and two seeds per configuration. Rate/native is $F_{\text{gen}}(d)/(hF_{\text{gen}}(10))$, not realized episode cost. The best observed count is 28/40 at $(1, 5)$; no monotonic denoising trend is assumed.

Allocation		Spatial task successes (each /10)					Nominal rate
h	d	Task 5	Task 6	Task 7	Task 8	Total	/ native
1	2	5	7	4	10	26/40	0.857
1	5	5	8	6	9	28/40	0.911
1	10	5	7	6	7	25/40	1.000
1	20	5	6	6	8	25/40	1.179
2	2	5	7	4	10	26/40	0.429
2	5	4	6	4	9	23/40	0.455
2	10	4	6	4	9	23/40	0.500
2	20	4	7	4	9	24/40	0.589
5	2	6	5	4	6	21/40	0.171
5	5	2	5	5	6	18/40	0.182
5	10	2	4	5	5	16/40	0.200
5	20	2	3	4	5	14/40	0.236

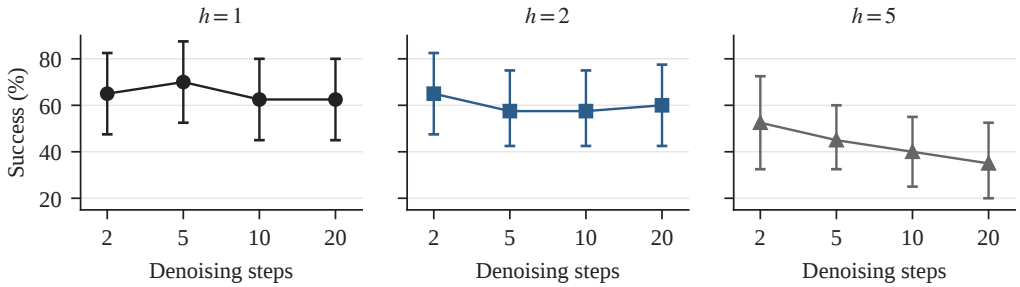


Figure 3: Discovery success by executed prefix h, with 40 attempts per configuration on four Spatial tasks. Connected estimates make the within-prefix response to denoising visible; whiskers are descriptive, task-conditional 95% paired state-bootstrap intervals. The d axes are categorical.

Calibration uses one real observation from Spatial task 0, initial state 0, and common explicit noise. For each $d \in \{2, 3, 5, 8, 10, 20\}$, profiling records solver timesteps and dense operation counts, separately from elapsed timing. Each setting receives 20 warm-up generations, followed by three randomized timing blocks of 20 measured calls per setting. The rollout evaluator also performs 20 warm-up generations once per configuration.

Host generation timing brackets prediction and completion synchronization; CUDA events measure the current-stream prediction interval. The outer rollout clock additionally includes observation processing, action handling, simulation/rendering, and record writing. Model loading, environment creation, reset, and warm-up are measured separately. Process sampling occurs approximately every five seconds; the largest recorded gap during discovery is below 5.03 seconds. Display processes remain on the measurement GPU.

Table 4: Discovery resource measurements. Generated-chunk totals use the declared dense-arithmetic boundary and charge partially executed final prefixes. Episode means include unsuccessful attempts. $(2, 2)$ has 26/40 observed successes at 42.9% of native arithmetic rate.

Allocation		Nominal / step	Mean per attempt		Success
h	d	GFLOPs	TFLOPs	Rollout (s)	Count
1	2	550.0	90.6	39.5	26/40
1	5	584.4	91.7	61.8	28/40
1	10	641.7	109.6	111.9	25/40
1	20	756.2	129.1	199.8	25/40
2	2	275.0	45.0	21.8	26/40
2	5	292.2	53.4	38.5	23/40
2	10	320.8	57.8	61.6	23/40
2	20	378.1	67.8	107.8	24/40
5	2	110.0	20.5	12.7	21/40
5	5	116.9	23.8	20.2	18/40
5	10	128.3	27.2	32.0	16/40
5	20	151.2	33.1	55.6	14/40

Calibration requires approximately 291 seconds and discovery 31,245 seconds (8.68 hours) of host elapsed time, including setup and bookkeeping. This search expenditure is additional to subsequent deployments. The experiment does not estimate an amortization threshold or a training–deployment cost crossover.

All measurements use an RTX 4060 Ti. Five-second process samples show no unrelated CUDA compute process during calibration, discovery, or the initially scheduled transfer phase; display processes remain. Process samples were not retained for priority completion or the beginning of continuation. Later continuation monitoring also shows no unrelated sampled CUDA compute process, but these coverage gaps make aggregate transfer rollout times descriptive. Host and CUDA-stream generation intervals are measured separately from the outer synchronous rollout interval. Warm-up is excluded from these rollout measurements.

A.1 Fixed-input generation timing

Table 5: Fixed-input full-generation calibration. Each setting has 60 calls in three randomized blocks. Range gives the minimum and maximum host block means, not a confidence interval. Host time falls more steeply with d than counted arithmetic. Host and CUDA columns average separate retained measurements of all 60 calls; the extra decimals distinguish their means, not timing accuracy.

d	GFLOPs	Host (ms)	CUDA (ms)	Host block range (ms)
2	550.0	217.891	217.862	214.5–220.1
3	561.5	269.990	269.970	265.3–272.6
5	584.4	374.321	374.300	366.8–378.1
8	618.8	529.746	529.726	519.9–536.1
10	641.7	629.624	629.604	623.2–642.2
20	756.2	1158.635	1158.608	1133.5–1172.6

The calibration stores full-generation intervals and global operator arithmetic counts, not separate vision, language-prefill, or action-expert timings and arithmetic totals. These records cannot attribute elapsed sensitivity to a particular component, memory traffic, omitted operation classes, synchronization, or software overhead. CUDA-event intervals are elapsed stream measurements, not summed active-kernel durations. No component-level decomposition is inferred.

A.2 Resource-coordinate sensitivity

Arithmetic savings and elapsed savings differ. Mean rollout time is 111.9 seconds per attempt for native deployment, 61.8 seconds for $(1, 5)$, and 21.8 seconds for $(2, 2)$. Their mean generated-chunk expenditures are 109.6, 91.7, and 45.0 TFLOPs, respectively. These are observed all-attempt means

under the same task allocation; they include different trajectories and outcomes. Table 4 reports all discovery configurations.

There is also an ordering reversal: $(1, 5)$ uses more arithmetic per attempt than $(2, 20)$ (91.7 versus 67.8 TFLOPs), but less rollout time (61.8 versus 107.8 seconds). These means include different trajectories. On a fixed input, reducing d from 10 to 5 lowers mean generation host time from 629.6 to 374.3 ms (40.5%), despite only 8.9% less counted arithmetic. The calibration therefore also separates arithmetic savings from elapsed savings without trajectory differences. It does not resolve their component-level cause. A FLOP-constrained choice is not necessarily a latency- or throughput-optimal choice.

As a descriptive sensitivity analysis, define calibrated host-time rate $b_T(h, d) = \overline{T}_{\mathrm{gen}}(d)/h$, the amortized fixed-input generation time per simulated control step, and normalize it by native host-time rate. For $(1, 2)$, $(2, 2)$, and $(5, 2)$, the normalized arithmetic rates are 0.857, 0.429, and 0.171, while the normalized host-time rates are 0.346, 0.173, and 0.069. Consequently, $(1, 2)$ is ineligible at a half-native arithmetic ceiling but eligible at a separately normalized half-native host-time ceiling. These resource constraints admit different refresh schedules; the choices do not establish quality optima. The frozen experiment remains arithmetic-based.

A.3 Language-shape transport audit

The profiled batch contains one sample, two $3 \times 360 \times 360$ image tensors, an eight-dimensional state tensor, and language-token and attention-mask tensors of length 20. The checkpoint exposes seven-dimensional actions and generates a 50-action chunk. Evaluation fixes batch size, image preprocessing, state/action dimensions, and chunk size. The tokenizer pads language only to the longest sequence in a batch. Reconstructing the declared task strings with the pinned tokenizer gives 15–17 tokens for discovery and 5–12 for transfer, versus 20 in calibration. We profile every distinct evaluated length for $d \in \{2, 5, 10\}$, preserving all other shapes, and test eligibility against the original absolute ceilings. Table 6 reports the range over evaluated lengths 5–17 against the original length-20 value.

Table 6: Dense generation arithmetic across evaluated language shapes. Budget tests retain the original absolute ceilings; none of the reported transfer configurations changes eligibility.

d	Evaluated lengths GFLOPs (range)	Length 20 GFLOPs	Eligibility changes
2	539.7–547.9	550.0	0
5	573.5–582.2	584.4	0
10	629.8–639.3	641.7	0

In particular, $(2, 10)$ remains within the original half-native ceiling for all evaluated lengths; shorter inputs reduce rather than increase its counted work. The audit changes neither policy execution nor the frozen selections.

A.4 Primary expenditure and runtime

The primary selected configuration $(2, 2)$ has a higher success count than $(2, 5)$ on seven tasks and ties on one (Table 9). There are 24 paired wins and seven losses. In a descriptive leave-one-task-out sensitivity analysis, the equally weighted difference over the remaining seven tasks stays positive, ranging from 9.3 to 12.1 percentage points. This post-outcome analysis does not extend inference to a new task population. The supplementary paired records retain all state and seed identities needed to reproduce the primary interval.

The retained provenance does not characterize numerical precision and attention execution for every model component. Table 8 therefore supports the stated model configuration, but not component- or kernel-level claims.

Table 7: Descriptive primary expenditure differences, selected minus comparator at $h = 2$, including all attempts. Intervals use the paired state bootstrap conditional on the eight tasks and were added after outcomes were observed. Timing intervals summarize sampling variation in the observed records; they do not account for unmonitored runtime interference during priority completion.

Mean per attempt	$d = 2$	$d = 5$	Difference	95% interval
Reference TFLOPs	41.1	48.7	-7.6	$[-10.6, -4.5]$
Observed rollout (s)	19.0	34.3	-15.2	$[-17.2, -13.3]$
Control steps	149.0	166.3	-17.3	$[-27.9, -6.5]$

Table 8: Runtime and evaluation conditions shared by all four confirmation execution phases. Module versions are the loaded versions reported at evaluation time. The supplement retains source digests and distinguishes conflicting installed distribution metadata from these loaded versions.

Setting	Value
Hardware	RTX 4060 Ti, 16 GB; driver 595.84
Runtime	Python 3.12.3; PyTorch 2.7.1+cu126; CUDA 12.6; NumPy 2.2.6; MuJoCo 3.8.1
Evaluation	Batch 1; synchronous simulation; nominal 20 Hz (50 ms control period)
Episode horizons	Spatial and Object: 280 control steps; Goal: 300
Precision controls	Autocast and TF32 disabled; deterministic algorithms enabled
Attention and reuse	Eager language/action attention with float32 scores; within-generation KV cache enabled
Compilation/warm-up	Model compilation disabled; 20 warm-up generations per condition
Evaluator provenance	Source revision and full digest retained in supplement

Table 9: Confirmation successes by task and configuration. Each task row has 20 attempts: ten paired initial states, each with two policy seeds. The total aggregates 160 attempts per configuration. The primary $(2, 2)$ setting exceeds $(2, 5)$ on seven tasks and ties on one.

	$h = 1$			$h = 2$ (primary pair)		$h = 5$	
Task	$d = 2$	$d = 5$	$d = 10$	$d = 2$	$d = 5$	$d = 2$	$d = 5$
Goal 5	14	15	14	16	14	18	16
Goal 6	18	17	17	18	15	19	18
Goal 7	20	20	20	19	18	18	18
Goal 8	19	18	18	18	14	18	15
Object 5	20	19	19	20	20	20	20
Object 6	19	18	19	20	17	18	18
Object 7	9	9	10	10	7	7	8
Object 8	18	20	18	18	17	20	20
Total (/160)	137	136	135	139	122	138	133

Table 10: Amended descriptive comparisons using all 160 attempts per setting. These contrasts were documented after collection began and are not originally prespecified or fully outcome-blind. Intervals are task-conditional, unadjusted, and do not test a quality-preservation margin.

Setting A	Setting B	A−B (pp)	95% interval (pp)
$(2, 2)$	$(1, 10)$	$+2.5$	$[-3.8, +8.1]$
$(2, 2)$	$(5, 2)$	$+0.6$	$[-6.3, +6.9]$
$(1, 2)$	$(2, 2)$	-1.3	$[-7.5, +5.0]$
$(1, 2)$	$(1, 10)$	$+1.3$	$[-5.0, +7.5]$

A.5 Complete confirmation outcomes

The four amended comparisons are reported in their documented order, using 20,000 paired state-bootstrap resamples and analysis seed 20260913. Table 1 gives their all-attempt resource means. In particular, the $(2, 2)$ minus $(5, 2)$ expenditure difference is $+25.0$ TFLOPs per attempt, with descriptive interval $[+22.8, +27.5]$. This interval describes generated-chunk arithmetic, including the different trajectory lengths; it is not an elapsed-time comparison or evidence of equivalent quality.

Measurement verification. The retained records support a complete internal reconciliation of planned identities, paired states, executed actions, and recorded outcomes. The checks below constrain recording and execution errors; they do not constitute an independent outcome replay in a separate simulator.

Table 11: Internal measurement-verification checklist.

Check	Evidence
Canonical coverage	1,120 confirmation attempts and 171,805 actions; the primary subset remains 320 attempts and 50,451 actions.
Pairing and dispatch	Planned identities, paired initial states and observations, generation counts, denoising counts, and applied actions reconcile.
Excluded records	Two repeated completed attempts are checked and excluded; one interrupted ten-action record has no completed outcome and is excluded.
Outcome path	All 1,122 raw completed records and 172,243 steps reconcile; generated actions match simulator actions and recorded outcome fields agree.
Shared predicate	Reward and success both use LIBERO's task-goal evaluation, traced to BDDL predicates, so their agreement is an internal consistency check.
Native parity	One Spatial episode matches the upstream path in initial state, observation, 70-action sequence, length, and outcome.
Remaining boundary	Independent, version-matched replay and direct terminal-state inspection remain outside this audit.

Analysis provenance. The three refinement comparisons, identities, and primary analysis were frozen before transfer outcomes. A deadline-driven amendment changed only execution order after 114 attempts per primary setting had been recorded; settings, identities, stopping rules, and inclusion remained fixed. Continuation completed the original identities, and deduplication retained the earliest completed record under the documented phase precedence. The frozen analyses use seed 20260912. Amended prefix contrasts and diagnostics use seed 20260913 and are descriptive because they were added after collection began. No response model was refitted on confirmation outcomes.

Analysis artifact. The accompanying supplement contains the primary episode and action records, planned identities, paired state values, calibrations, frozen predictions, runtime provenance, and analysis code. Its bundled verification procedure reproduces the primary counts, interval, task table, and descriptive expenditure analysis. Documentation specifies the selector features, penalized objective, candidate set, and normalization. Checkpoint weights and simulator assets are not bundled. The continuing repository additionally retains a completed confirmation snapshot and verification code that reproduce the full seven-setting resource table, original and amended comparisons, and task-level counts. The submitted supplement remains separate from these post-submission additions.

B Frozen Selection and Diagnostic Analyses

We compare two logistic response-model families. A compute-only model uses a quadratic in log normalized rate. An allocation model uses $\log h$, $\log(d/10)$, both squared terms, and their interaction. Each family is fitted with ridge penalty 1 or 10. Leave-one-task-out log loss selects the model, which is then refitted on the discovery set. No individual seed repeat is held out while another repeat from the same task remains in training.

The candidate set is $h \in \{1, 2, 3, 5\}$ crossed with $d \in \{2, 3, 5, 8, 10, 20\}$: the 12 discovery configurations and 12 untested configurations. Their calibrated costs determine eligibility; the fitted model

Table 12: Discovery-based selector choices. The learned choices and comparisons were frozen before transfer; additional selector baselines were documented after collection began. The empirical selector uses the measured grid; other selectors use the declared candidate set. The simple rule exactly matches the learned choices. Selected rate/native is the learned recommendation's arithmetic rate. Frozen comparators and predicted gaps appear in Table 13; the half-native row is primary.

	Selector choice (h, d)				Learned
Ceiling/native	Learned	Simple	Empirical	Compute-only	rate/native
0.25	$(5, 2)$	$(5, 2)$	$(5, 2)$	$(5, 20)$	0.171
0.50 (primary)	$(2, 2)$	$(2, 2)$	$(2, 2)$	$(2, 10)$	0.429
1.00	$(1, 2)$	$(1, 2)$	$(1, 5)$	$(1, 10)$	0.857

Table 13: All original frozen denoising tests: selected minus comparator, holding h fixed within each row. Predicted and observed gaps and intervals are percentage points. The half-native row is primary; secondary intervals are descriptive, task-conditional, and unadjusted.

Ceiling	Setting (h, d)		Success difference (pp)		
/ native	Selected	Comparator	Predicted	Observed	95% interval
0.25	$(5, 2)$	$(5, 5)$	$+6.5$	$+3.1$	$[-2.5, +8.8]$
0.50 (primary)	$(2, 2)$	$(2, 5)$	$+4.0$	$+10.6$	$[+4.4, +16.9]$
1.00	$(1, 2)$	$(1, 5)$	$+2.2$	$+0.6$	$[-5.0, +6.3]$

selects the highest predicted success under each ceiling. Ties favor lower compute, then lexicographic order.

The simple rule minimizes denoising count, then chooses the shortest feasible prefix. It reproduces every learned recommendation (Table 12); transfer outcomes cannot distinguish these selectors at the tested budgets. The compute-only choices use the original ridge-10 fit. Its quarter- and half-budget choices are absent from the transfer plan, so this evaluation cannot compare those decisions on transfer tasks.

The selected allocation model has log loss 0.7048, compared with 0.7067 for the best compute-only model. This small difference provides no evidence of an added predictive advantage. Under the same mean natural-log binary cross-entropy, constant $p = 0.5$ scores 0.6931 and an intercept-only model fitted separately on each training fold scores 0.7170. The response-model probabilities should therefore not be read as calibrated task-transfer forecasts. The primary test evaluates a recommendation shared with simpler selectors, so it does not establish an additional benefit from model-based ranking.