Refinement versus Refresh: Compute Allocation for Closed-Loop Robot Policies

Anonymous Author(s)
Affiliation
Address
email

Abstract

Chunked robot policies expose two distinct uses of inference compute: refining the current action prediction and obtaining a new prediction from a fresher observation. We study how to allocate compute between these operations in a frozen SmolVLA policy. We formalize deployment compute rate as $b(h, d) = F_{\text{gen}}(d)/h$, where d controls denoising refinement and h is the number of predicted actions executed before replanning. Full-generation profiling reveals a substantial denoising-independent cost per generation: 82.1% of native counted arithmetic is independent of denoising, so reducing refinement alone cannot proportionally reduce deployment compute and tighter budgets can change which refresh schedules are feasible. We evaluate this trade-off with a paired 480-attempt, 12-configuration LIBERO discovery study and a completed 1,120-attempt confirmation study on eight tasks excluded from selection. At the primary two-action prefix, reducing denoising from five steps to two increases success from 122/160 to 139/160, a paired difference of 10.6 percentage points with a task-conditional, descriptive 95% interval of 4.4–16.9 points, while reducing reference-calibrated arithmetic by 5.9% per control step. At two denoising steps, a five-action prefix records 138/160 successes while using 60.7% less arithmetic per attempt than the two-action prefix, although the quality difference remains unresolved. A simple cost-aware rule matches the fitted response model's choices at three compute ceilings, so fitting provides no demonstrated selection advantage. These results show that more inference refinement need not improve closed-loop control and narrow the broader allocation question: when is the fixed cost of obtaining a fresh prediction worth paying, and how much refinement should that prediction receive?

1 Introduction

A robot policy can reduce inference work in two ways: spend less computation refining each action prediction, or execute more actions before requesting a new prediction. These interventions affect behavior differently. Reduced refinement changes the prediction itself; longer execution prefixes delay observation-conditioned revision and expose later positions of the predicted action sequence. Their relative value should therefore be evaluated through closed-loop outcomes under a common resource constraint.

We ask: *given a fixed inference budget, how should a robot divide computation between prediction refinement and observation-conditioned refresh?* The cost of a fresh prediction makes this question nontrivial. In our frozen SmolVLA implementation [1], work independent of denoising accounts for 82.1% of native counted arithmetic. Reducing refinement therefore leaves a substantial cost per generation: sufficiently tight budgets require less frequent refresh even at low denoising counts. This cost boundary determines what is feasible, but not which feasible setting controls the robot best.

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

We separate two empirical questions. First, does additional denoising improve success at a fixed execution prefix? Second, does more frequent refresh justify its additional expenditure? We evaluate twelve configurations in a paired 480-attempt discovery study, then seven configurations in 1,120 confirmation attempts on eight tasks excluded from selection. Three frozen comparisons test the first question. Later-amended descriptive prefix comparisons address the second; they neither isolate observation freshness from action-position effects nor establish an optimal allocation rule.

At the primary two-action prefix, two denoising steps yield 139/160 successes versus 122/160 for five steps, with lower reference arithmetic. At two denoising steps, however, a five-action prefix yields 138/160 successes at 60.7% less arithmetic per attempt than the two-action prefix; their quality difference remains unresolved. A simple cost-aware rule also matches the fitted response model's choices at all three ceilings, with no demonstrated selection advantage from fitting.

The contributions are:

- a refinement–refresh allocation formulation and full-generation calibration exposing the fixed cost of obtaining a fresh prediction;
- a paired twelve-configuration study reporting both nominal compute rate and realized episode expenditure; and
- a completed task-transfer evaluation separating frozen refinement tests from descriptive refresh comparisons, with their uncertainty and selection provenance retained.

2 Related Work

Resource-allocation studies motivate evaluating configurations across explicit budgets. Chinchilla studies model-size and training-data allocation [2]; embodied scaling studies consider world modeling, behavior cloning, and closed-loop forecasting and planning [3, 4]. Our comparison instead varies deployment refinement and execution prefix within a fixed policy. Work distinguishing FLOP-optimal from elapsed- or monetary-cost-optimal allocation [5] also motivates retaining separate resource coordinates.

VLA-Perf characterizes inference performance across denoising, chunk size, asynchronous execution, and hardware/runtime choices [6]. SmolVLA studies asynchronous inference and robot behavior [1]. Here predicted chunk size remains fixed, while executed prefix changes; task success is measured alongside computation. RoboMonkey studies test-time sampling and action verification [7], an adjacent allocation axis that would require counting candidate generation and verification costs.

Adaptive VLA scheduling is directly related. DVAC uses denoising-derived action variance to choose an execution prefix before replanning [8]. ElegantVLA schedules vision–language recomputation or reuse together with action denoising modes [9]. Our study instead compares fixed configurations under explicit arithmetic ceilings and freezes separate-task denoising comparisons. It neither introduces an adaptive scheduler nor establishes an advantage over these methods; their reported speedups have different policies, hardware, and measurement boundaries.

3 Refinement–Refresh Allocation

3.1 Compute rate and episode expenditure

Let d be the denoising count and h the number of successive predicted actions executed before replanning. Let $F_{\mathrm{gen}}(d)$ denote reference-calibrated neural arithmetic for a full action-chunk generation at denoising setting d. For complete execution prefixes, the nominal compute rate per simulated control step is

$$b(h, d) = \frac{F_{\mathrm{gen}}(d)}{h}.$$

For a candidate set $\mathcal{C}$ and budget ceiling B, the allocation target is

$$(h^\star, d^\star)(B) \in \arg \max_{(h,d)\in\mathcal{C}:\, b(h,d)\le B} p_{\mathrm{success}}(h, d). \tag{1}$$

A configuration need not exhaust the ceiling. This is a deployment compute-rate constraint for a fixed checkpoint, rather than a training-plus-deployment objective.

We distinguish this rate from total arithmetic per attempted episode. If attempt e executes T_e actions, its generated-chunk expenditure is

$$C_e = \lceil T_e/h \rceil \, F_{\text{gen}}(d). \tag{2}$$

The final generated chunk is charged even if only part of its execution prefix is used. All-attempt means include unsuccessful episodes. Different trajectories and termination times can therefore change expenditure at the same nominal rate. FLOPs, host elapsed time, and CUDA-stream time retain separate units and measurement boundaries.

3.2 The fixed cost of a fresh prediction

Complete-generation profiling counts matrix products, convolutions, and attention across vision, language prefill, and the action expert, using two FLOPs per multiply-add. Embedding lookups, normalization, activations, elementwise arithmetic, and external preprocessing are excluded. We therefore refer to *dense neural arithmetic*, not an exact count of every scalar operation. Calibration uses one real development observation and fixed tensor shapes; its protocol is detailed in Appendix A.

For the reference input, the measured arithmetic satisfies, to rounding,

$$F_{\text{gen}}(d) \simeq 527.108 + 11.457d \quad \text{GFLOPs}$$

at all six calibrated denoising settings. Native generation at $d = 10$ costs 641.675 GFLOPs, of which 82.1% is independent of d. Halving denoising from ten to five saves 8.9% of counted generation work; reducing it to two saves 14.3%. Equivalently, relative to native deployment $(h, d) = (1, 10)$,

$$\frac{b(h,d)}{b_{\text{native}}} \simeq \frac{0.821 + 0.179(d/10)}{h}.$$

We prespecify compute-rate ceilings of 0.25, 0.50, and 1.00 times native. Eligibility uses measured per-d values, not the rounded affine expression. For example, $(1, 2)$ costs 0.857 times native rate and is ineligible at half-native budget, whereas $(2, 2)$ costs 0.429 times native. The denoising-independent work alone excludes $h = 1$ at half-native budget and $h \leq 2$ at quarter-native budget. A tighter budget can therefore force a longer execution prefix before any quality-based choice is made. This is an implementation-specific arithmetic floor under the declared counting boundary, not a universal VLA scaling law.

Calibration uses a fixed input shape; a separate language-length audit leaves eligibility unchanged under the original absolute ceilings (Appendix A.3). Reference arithmetic remains distinct from task-shaped arithmetic and elapsed time. The cost model defines feasible schedules; the following experiments ask which feasible expenditures help closed-loop control.

4 Controlled Study and Frozen Comparisons

4.1 Discovery and confirmation protocol

We evaluate the public `HuggingFaceVLA/smolvla_libero` checkpoint on LIBERO [10]. Predicted chunk length remains 50 actions; native execution uses $(h, d) = (1, 10)$. We hold weights, sensors, preprocessing, action normalization, control period, and task-specific episode horizons fixed.

Discovery crosses $h \in \{1, 2, 5\}$ and $d \in \{2, 5, 10, 20\}$ on Spatial tasks 5–8, five initial states per task, and two policy seeds. Each of the 40 paired task–state–seed identities receives all twelve configurations, giving 480 attempts. Tasks 5 and 6 overlap the pilot; tasks 7 and 8 were selected before observing their outcomes in this study. Condition order is randomized within each paired identity.

Discovery is used to freeze three budget-specific denoising comparisons before confirmation. Their selected settings and comparators, together with native, give seven configurations. Confirmation evaluates all seven on Goal and Object tasks 5–8, ten states per task, and two seeds: 160 attempts per configuration, totaling 1,120. These eight tasks are excluded from configuration selection; prior exposure during unrelated development is not assessed. The sample allocation prioritizes task and state coverage; no quality-preservation margin is prespecified. Appendix A.5 distinguishes the original tests from the later-amended descriptive comparisons.

Initial simulator states and observations match across conditions. Policy noise is keyed by task, state, seed, and simulation tick, preserving a common noise key whenever conditions generate at the same tick. Solver timesteps and a constant-field integration check verify that every d traverses the full interval from 1 to 0; fewer steps change the discretization, not the integration endpoint. Action and outcome journals reconcile independently across software boundaries (Appendix A).

Success averages both seeds within each state, then states and tasks with equal task weights. Descriptive 95% intervals use 20,000 bootstrap draws of states within the fixed tasks, retaining paired seeds and configurations. They condition on four discovery or eight confirmation tasks; secondary intervals are unadjusted. Arithmetic includes every generated chunk, including failures and partly consumed final prefixes. Timing measurements and their monitoring limits are reported separately in Appendix A.1.

4.2 Discovery: quality does not increase monotonically with refinement

Native $(1, 10)$ succeeds in 25/40 attempts. The highest observed count is 28/40 for $(1, 5)$; $(1, 2)$ and $(2, 2)$ each succeed in 26/40. At $h = 5$, successes are 21, 18, 16, and 14 out of 40 as d increases from 2 to 20. Figure 1(a) shows every configuration. These point estimates do not justify imposing monotonic improvement with denoising. The small paired sample also does not certify equal quality for configurations with similar observed counts.

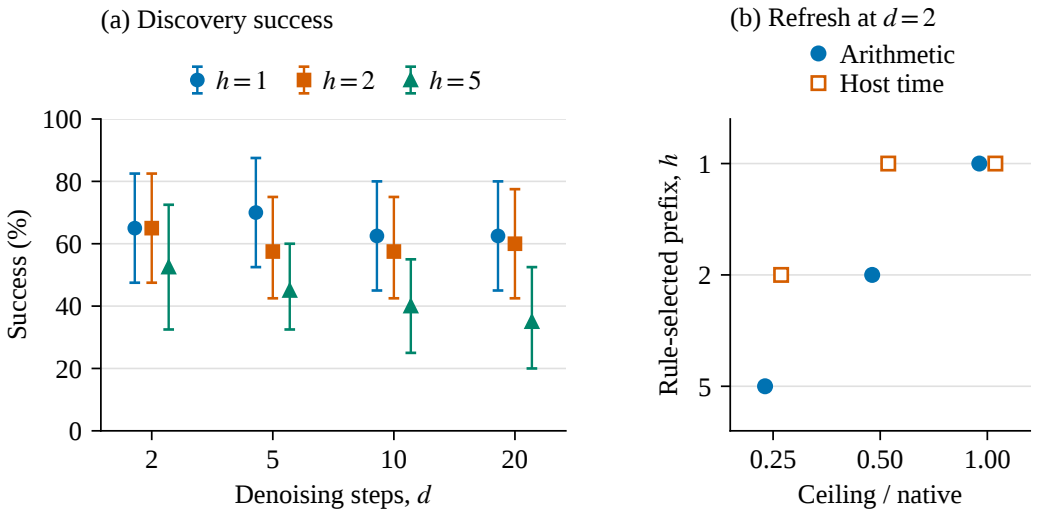


Figure 1: Quality and feasibility are distinct. (a) Discovery success, with 40 attempts per configuration. Colors and shapes identify executed prefix h; bars give descriptive, task-conditional 95% paired state-bootstrap intervals. The d axis is categorical; horizontal offsets only separate overlapping bars. (b) A minimum-denoising, shortest-feasible-prefix rule selects different prefixes under separately normalized arithmetic and fixed-input host-time ceilings. The h axis shows the selected prefix categories, not elapsed time. These are feasibility-dependent choices, not established quality optima; timing details are in Appendix A.2.

The resource coordinate also matters: nominal rate differs from realized episode expenditure, and arithmetic savings differ from elapsed savings. Figure 1(b) illustrates how substituting a host-time constraint changes feasible refresh choices. This sensitivity analysis does not change the frozen arithmetic-based experiment. Full task and resource tables appear in Appendix A.

4.3 Frozen tests and the role of selection

A response model fitted only to discovery selects $(5, 2)$, $(2, 2)$, and $(1, 2)$ at the quarter-, half-, and full-native ceilings. Their frozen comparators are $(5, 5)$, $(2, 5)$, and $(1, 5)$, respectively; the half-native comparison is primary. All selected settings are existing discovery grid points, so confirmation tests task transfer rather than extrapolation to an untested configuration. Each pair changes denoising

at a fixed prefix. These are comparisons under common ceilings, not equal-rate tests of how to divide compute between refinement and refresh.

At the primary ceiling, $(2, 2)$ has the highest eligible discovery success. Its comparator $(2, 5)$ ties with $(2, 10)$ for the highest count among distinct eligible alternatives and wins the lower-cost tie-break. A simple rule choosing minimum denoising and then the shortest feasible prefix makes the same choices as the fitted model at all three ceilings. This baseline was documented after collection began and is interpreted descriptively. Transfer outcomes cannot demonstrate a selection advantage between rules that make identical choices. Model fitting, the full candidate set, and selector diagnostics are retained in Appendix B.

5 Refinement and Refresh on Transfer Tasks

5.1 Complete confirmation overview

All 1,120 originally planned confirmation attempts are complete. Table 1 reports the seven configurations on the same paired identities. The audit excludes two repeated completed records under the documented precedence rule; continuation completes the original design, not an independent replication (Appendix A). We first report the frozen denoising tests, then the amended prefix comparisons.

Table 1: Confirmation overview: 160 attempts per configuration on eight tasks. h is the executed prefix and d the denoising count. Nominal rate/native is $F_{\text{gen}}(d)/(hF_{\text{gen}}(10))$. All-attempt arithmetic uses the length-20 reference calibration and charges every generation. The primary pair is $(2, 2)$ versus $(2, 5)$.

Allocation		Success		Nominal rate	Mean per attempt	
h	d	Count	(%)	/ native	TFLOPs	Steps
1	2	137/160	85.6	0.857	83.3	151.4
1	5	136/160	85.0	0.911	89.1	152.5
1	10	135/160	84.4	1.000	98.4	153.4
2	2	139/160	86.9	0.429	41.1	149.0
2	5	122/160	76.3	0.455	48.7	166.3
5	2	138/160	86.3	0.171	16.1	144.9
5	5	133/160	83.1	0.182	18.5	156.2

5.2 Refinement: the frozen denoising comparisons

At the primary two-action prefix, two denoising steps yield 139/160 successes versus 122/160 for five steps. The paired difference is $+10.6$ percentage points, with descriptive 95% interval $[+4.4, +16.9]$, conditional on the eight tasks. The selected setting also uses 5.9% less nominal reference arithmetic per control step: 275.0 versus 292.2 GFLOPs. Thus reducing refinement improves observed success while reducing counted work at this fixed prefix.

Mean reference arithmetic per attempt falls from 48.7 to 41.1 TFLOPs (15.6%). This larger reduction includes changed trajectories and termination: mean episode length is 149.0 versus 166.3 control steps. All failures reach the task horizon (21 versus 38); the remaining attempts terminate successfully. The primary direction is positive on seven tasks and tied on one, with 24 paired wins and seven losses. Appendix A.4 gives task sensitivity and expenditure intervals. These measurements do not establish a controlled physical-latency improvement. Figure 2 places the primary effect alongside all other frozen and amended comparisons.

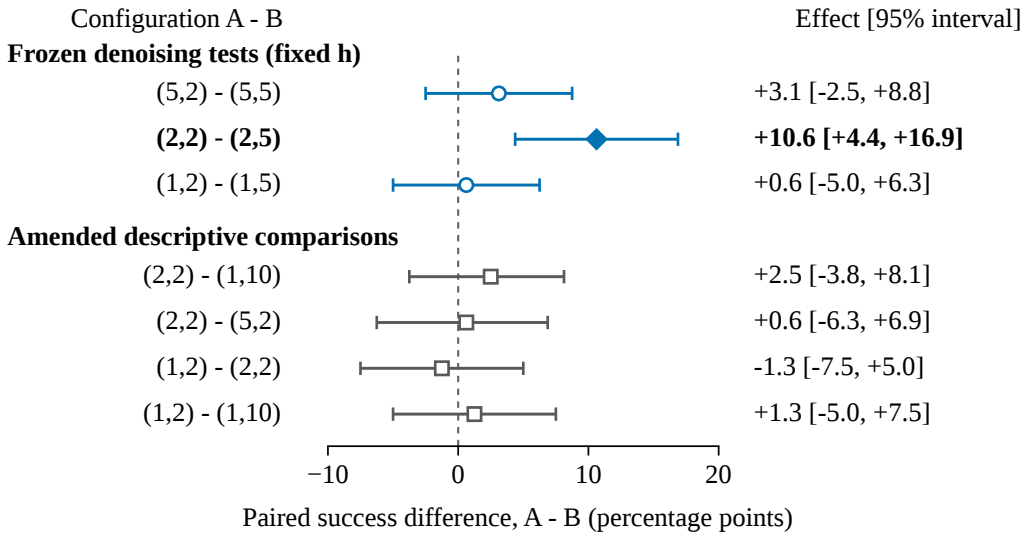


Figure 2: Confirmation effects: 160 paired attempts per setting on eight tasks. Rows compare (h, d) as A minus B; bars are task-conditional, descriptive 95% paired state-bootstrap intervals (secondary intervals unadjusted). The filled diamond marks the prespecified primary test. All three frozen tests and all four amended contrasts appear in their documented order; the latter were not fully outcome-blind. Intervals crossing zero do not establish equivalence. Predicted and observed gaps are tabulated in Tables 11 and 9.

The other two frozen tests have positive point estimates but intervals spanning zero (Figure 2). At $h = 5$, two versus five denoising steps give 138/160 versus 133/160 successes; at $h = 1$, they give 137/160 versus 136/160. The largest observed benefit is at the primary prefix. The primary gap exceeds its frozen prediction, while matching the signs of three point estimates does not establish calibrated effect-size prediction. These results do not demonstrate a uniform denoising benefit across prefixes.

5.3 Refresh: descriptive prefix comparisons

We now hold denoising fixed and compare expenditures on different refresh schedules. These contrasts were documented after collection began, when individual outcomes were visible; they are not originally prespecified or fully outcome-blind (Appendix A.5).

At $d = 2$, $(2, 2)$ exceeds $(5, 2)$ by just one success: 139/160 versus 138/160, a paired difference of $+0.6$ percentage points $[-6.3, +6.9]$. Yet $(5, 2)$ uses 60% less nominal arithmetic per control step and 60.7% less reference arithmetic per attempt: 16.1 versus 41.1 TFLOPs. The interval permits meaningful differences in either direction; similar counts do not certify equivalent quality or justify a non-inferiority claim.

Refreshing every action with $(1, 2)$ instead of every two actions gives two fewer successes, a difference of -1.3 points $[-7.5, +5.0]$, at 83.3 versus 41.1 TFLOPs per attempt. More frequent refresh has no demonstrated quality benefit in these contrasts. Native-relative comparisons, which also change denoising, are reported separately in Appendix A.5.

The discovery point estimates favor shorter prefixes at $d = 2$ (Section 4.2), but their quality advantage is not established on transfer. The task populations differ, so this comparison does not identify a causal task effect. The cost model explains feasibility, but does not determine the quality ranking on new tasks.

6 Discussion and Conclusion

Full-generation cost determines which refresh schedules are feasible, while closed-loop outcomes determine which feasible expenditure is useful. The fixed cost of a fresh prediction makes refinement reduction alone insufficient at tight arithmetic ceilings. Within those constraints, the frozen primary result favors less denoising at a two-action prefix. The descriptive prefix comparisons identify a much cheaper five-action setting whose quality difference remains unresolved. These findings constrain the allocation question without identifying a universal optimum, and the fitted selector adds no demonstrated advantage over the simple rule.

Interpretation and limits. The two deployment controls are not isolated causal interventions on prediction quality and observation freshness. Changing denoising changes numerical integration; changing prefix changes both observation age and executed action positions. The mechanism behind the benefit of coarser sampling remains open. The study covers one checkpoint and twelve fixed tasks in simulation, with only four discovery task-level validation folds. Intervals are task-conditional, secondary contrasts are unadjusted, and no quality-preservation margin was prespecified. Two steps is the minimum denoising tested; $d = 1$ is untested. The arithmetic boundary excludes several operation classes; language lengths are audited but image shape is fixed. Sampled GPU checks have coverage gaps, and simulated control steps do not establish real-time responsiveness.

Next steps. A matched stale-input diagnostic on the 40 discovery identities changes the supplied images and proprioception refresh from every tick to every five ticks while generating at every tick at $(h, d) = (1, 5)$. Success falls from 28/40 with fresh inputs to 20/40 with held inputs, a paired difference of -20.0 percentage points with a descriptive task-conditional 95% interval $[-32.5, -7.5]$. This supports an observation-freshness contribution, but the step journal logs action age rather than per-step supplied-input hashes or observation age; the intervention therefore does not isolate every implementation detail. Completing the input instrumentation and checking action positions can further separate the mechanisms behind prefix effects. A new, predeclared quality-preservation test for the cheaper $(5, 2)$ setting and replication on another policy are needed to assess whether its observed cost advantage supports deployment. The broader objective is to determine when the fixed cost of a fresh prediction is worth paying, and how much refinement that prediction needs.

References

[1] Mustafa Shukor, Dana Aubakirova, Francesco Capuano, Pepijn Kooijmans, Steven Palma, Adil Zouitine, Michel Aractingi, Caroline Pascal, Martino Russi, Andres Marafioti, Simon Alibert, Matthieu Cord, Thomas Wolf, and Remi Cadene. SmolVLA: A vision-language-action model for affordable and efficient robotics. arXiv preprint arXiv:2506.01844v1, 2025. URL https://arxiv.org/abs/2506.01844v1.

[2] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Thomas Hennigan, Eric Noland, Katherine Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karén Simonyan, Erich Elsen, Oriol Vinyals, Jack Rae, and Laurent Sifre. An empirical analysis of compute-optimal large language model training. In *Advances in Neural Information Processing Systems*, volume 35, 2022. doi: 10.52202/068431-2176. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/c1e2faff6f588870935f114ebe04a3e5-Abstract.html.

[3] Tim Pearce, Tabish Rashid, David Bignell, Raluca Georgescu, Sam Devlin, and Katja Hofmann. Scaling laws for pre-training agents and world models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48542–48562. PMLR, 2025. URL https://proceedings.mlr.press/v267/pearce25a.html.

[4] Mustafa Baniodeh, Kratarth Goel, Scott Ettinger, Carlos Fuertes, Ari Seff, Tim Shen, Cole Gulino, Chenjie Yang, Ghassen Jerfel, Dokook Choe, Rui Wang, Benjamin Charrow, Vinutha Kallem, Sergio Casas, Rami Al-Rfou, Benjamin Sapp, and Dragomir Anguelov. Scaling laws of motion forecasting and planning – technical report. arXiv preprint arXiv:2506.08228v2, 2025. URL https://arxiv.org/abs/2506.08228v2.

[5] Nikhil Sardana, Jacob Portes, Sasha Doubov, and Jonathan Frankle. Beyond chinchilla-optimal: Accounting for inference in language model scaling laws. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 43445–43460. PMLR, 2024. URL https://proceedings.mlr.press/v235/sardana24a.html.

[6] Wenqi Jiang, Jason Clemons, Karu Sankaralingam, and Christos Kozyrakis. How fast can I run my VLA? demystifying VLA inference performance with VLA-Perf. arXiv preprint arXiv:2602.18397v1, 2026. URL https://arxiv.org/abs/2602.18397v1.

[7] Jacky Kwok, Christopher Agia, Rohan Sinha, Matt Foutter, Shulu Li, Ion Stoica, Azalia Mirhoseini, and Marco Pavone. RoboMonkey: Scaling test-time sampling and verification for vision-language-action models. In *Proceedings of the 9th Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 3200–3217. PMLR, 2025. URL https://proceedings.mlr.press/v305/kwok25a.html.

[8] Xiangdong Feng, Yuxuan Cheng, Chen Shi, Boyao Han, Yuxuan Yan, Yitong Hong, Zhuotao Tian, and Li Jiang. Denoising tells when to replan: Denoising-variance adaptive chunking for flow-based robot policies. arXiv preprint arXiv:2606.03847v1, 2026. URL https://arxiv.org/abs/2606.03847v1.

[9] Ye Li, Huanan Liu, Kangye Ji, Yuan Meng, Jiajun Fan, Yuansong Wang, Shiyu Qin, Chenglei Wu, Shu-Tao Xia, and Zhi Wang. ElegantVLA: Learning when to think for efficient vision-language-action models. arXiv preprint arXiv:2605.29438v1, 2026. URL https://arxiv.org/abs/2605.29438v1.

[10] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. LIBERO: Benchmarking knowledge transfer for lifelong robot learning. In *Advances in Neural Information Processing Systems*, volume 36, 2023. doi: 10.52202/075280-1939. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets_and_Benchmarks.html.

A Detailed Outcomes and Resource Measurements

The checkpoint is `HuggingFaceVLA/smolvla_libero`, revision `6721902bc4d61e50a3bfdb11dfb4cb626f05d102`. Predicted chunk size is 50; native execution uses $h = 1, d = 10$. Discovery uses Spatial tasks 5–8, initial states 10–14, and policy seeds 2000 and 2001. The 20 task–state blocks contain both seeds and all 12 configurations. Confirmation uses Goal and Object tasks 5–8, states 20–29, and policy seeds 3000 and 3001. Its 80 task–state blocks each contain both seeds and all seven configurations. Discovery bootstrap resampling uses 20,000 draws with analysis seed 20260913.

Table 2: Discovery successes on Spatial tasks 5–8. Each task contributes five states and two seeds per configuration. Rate/native is $F_{\text{gen}}(d)/(hF_{\text{gen}}(10))$, not realized episode cost. The best observed count is 28/40 at $(1, 5)$; no monotonic denoising trend is assumed.

Allocation		Spatial task successes (each /10)					Nominal rate
h	d	Task 5	Task 6	Task 7	Task 8	Total	/ native
1	2	5	7	4	10	26/40	0.857
1	5	5	8	6	9	28/40	0.911
1	10	5	7	6	7	25/40	1.000
1	20	5	6	6	8	25/40	1.179
2	2	5	7	4	10	26/40	0.429
2	5	4	6	4	9	23/40	0.455
2	10	4	6	4	9	23/40	0.500
2	20	4	7	4	9	24/40	0.589
5	2	6	5	4	6	21/40	0.171
5	5	2	5	5	6	18/40	0.182
5	10	2	4	5	5	16/40	0.200
5	20	2	3	4	5	14/40	0.236

Calibration uses one real observation from Spatial task 0, initial state 0, and common explicit noise. For each $d \in \{2, 3, 5, 8, 10, 20\}$, profiling records solver timesteps and dense operation counts, separately from elapsed timing. Each setting receives 20 warm-up generations, followed by three randomized timing blocks of 20 measured calls per setting. The rollout evaluator also performs 20 warm-up generations once per configuration.

Host generation timing brackets prediction and completion synchronization; CUDA events measure the current-stream prediction interval. The outer rollout clock additionally includes observation processing, action handling, simulation/rendering, and journal writes. Model loading, environment creation, reset, and warm-up are measured separately. Process sampling occurs approximately every five seconds; the largest recorded gap during discovery is below 5.03 seconds. Display processes remain on the measurement GPU.

Calibration requires approximately 291 seconds and discovery 31,245 seconds (8.68 hours) of host elapsed time, including setup and bookkeeping. This search expenditure is additional to subsequent deployments. The experiment does not estimate an amortization threshold or a training–deployment cost crossover.

All measurements use an RTX 4060 Ti. Five-second process samples show no unrelated CUDA compute process during calibration, discovery, or the initially scheduled transfer phase; display processes remain. Process samples were not retained for priority completion or the beginning of continuation. Later continuation monitoring also shows no unrelated sampled CUDA compute process, but these coverage gaps make aggregate transfer rollout times descriptive. Host and CUDA-stream generation intervals are measured separately from the outer synchronous rollout interval. Warm-up is excluded from these rollout measurements.

A.1 Fixed-input generation timing

The calibration stores full-generation intervals and global operator arithmetic counts, not separate vision, language-prefill, or action-expert timings and arithmetic totals. These records cannot attribute elapsed sensitivity to a particular component, memory traffic, omitted operation classes, synchroniza-

Table 3: Discovery resource measurements. Generated-chunk totals use the declared dense-arithmetic boundary and charge partially executed final prefixes. Episode means include unsuccessful attempts. $(2, 2)$ has 26/40 observed successes at 42.9% of native arithmetic rate.

Allocation		Nominal / step	Mean per attempt		Success
h	d	GFLOPs	TFLOPs	Rollout (s)	Count
1	2	550.0	90.6	39.5	26/40
1	5	584.4	91.7	61.8	28/40
1	10	641.7	109.6	111.9	25/40
1	20	756.2	129.1	199.8	25/40
2	2	275.0	45.0	21.8	26/40
2	5	292.2	53.4	38.5	23/40
2	10	320.8	57.8	61.6	23/40
2	20	378.1	67.8	107.8	24/40
5	2	110.0	20.5	12.7	21/40
5	5	116.9	23.8	20.2	18/40
5	10	128.3	27.2	32.0	16/40
5	20	151.2	33.1	55.6	14/40

Table 4: Fixed-input full-generation calibration. Each setting has 60 calls in three randomized blocks. Range gives the minimum and maximum host block means, not a confidence interval. Host time falls more steeply with d than counted arithmetic. Host and CUDA columns average separate retained measurements of all 60 calls; the extra decimals distinguish their means, not timing accuracy.

d	GFLOPs	Host (ms)	CUDA (ms)	Host block range (ms)
2	550.0	217.891	217.862	214.5–220.1
3	561.5	269.990	269.970	265.3–272.6
5	584.4	374.321	374.300	366.8–378.1
8	618.8	529.746	529.726	519.9–536.1
10	641.7	629.624	629.604	623.2–642.2
20	756.2	1158.635	1158.608	1133.5–1172.6

tion, or software overhead. CUDA-event intervals are elapsed stream measurements, not summed active-kernel durations. No component-level decomposition is inferred.

A.2 Resource-coordinate sensitivity

Arithmetic savings and elapsed savings differ. Mean rollout time is 111.9 seconds per attempt for native deployment, 61.8 seconds for $(1, 5)$, and 21.8 seconds for $(2, 2)$. Their mean generated-chunk expenditures are 109.6, 91.7, and 45.0 TFLOPs, respectively. These are observed all-attempt means under the same task allocation; they include different trajectories and outcomes. Table 3 reports all discovery configurations.

There is also an ordering reversal: $(1, 5)$ uses more arithmetic per attempt than $(2, 20)$ (91.7 versus 67.8 TFLOPs), but less rollout time (61.8 versus 107.8 seconds). These means include different trajectories. On a fixed input, reducing d from 10 to 5 lowers mean generation host time from 629.6 to 374.3 ms (40.5%), despite only 8.9% less counted arithmetic. The calibration therefore also separates arithmetic savings from elapsed savings without trajectory differences. It does not resolve their component-level cause. A FLOP-constrained choice is not necessarily a latency- or throughput-optimal choice.

As a descriptive sensitivity analysis, define calibrated host-time rate $b_T(h, d) = \overline{T}_{\mathrm{gen}}(d)/h$, the amortized fixed-input generation time per simulated control step, and normalize it by native host-time rate. For $(1, 2)$, $(2, 2)$, and $(5, 2)$, the normalized arithmetic rates are 0.857, 0.429, and 0.171, while the normalized host-time rates are 0.346, 0.173, and 0.069. Consequently, $(1, 2)$ is ineligible at a half-native arithmetic ceiling but eligible at a separately normalized half-native host-time ceiling. Figure 1(b) shows the resulting choices at all three ceilings. These resource constraints admit different refresh schedules; the choices do not establish quality optima. The frozen experiment remains arithmetic-based.

A.3 Language-shape transport audit

The profiled batch contains one sample, two $3 \times 360 \times 360$ image tensors, an eight-dimensional state tensor, and language-token and attention-mask tensors of length 20. The checkpoint exposes seven-dimensional actions and generates a 50-action chunk. Evaluation fixes batch size, image preprocessing, state/action dimensions, and chunk size. The tokenizer pads language only to the longest sequence in a batch. Reconstructing the declared task strings with the pinned tokenizer gives 15–17 tokens for discovery and 5–12 for transfer, versus 20 in calibration. We profile every distinct evaluated length for $d \in \{2, 5, 10\}$, preserving all other shapes, and test eligibility against the original absolute ceilings. Table 5 reports the range over evaluated lengths 5–17 against the original length-20 value.

Table 5: Dense generation arithmetic across evaluated language shapes. Budget tests retain the original absolute ceilings; none of the reported transfer configurations changes eligibility.

d	Evaluated lengths GFLOPs (range)	Length 20 GFLOPs	Eligibility changes
2	539.7–547.9	550.0	0
5	573.5–582.2	584.4	0
10	629.8–639.3	641.7	0

In particular, $(2, 10)$ remains within the original half-native ceiling for all evaluated lengths; shorter inputs reduce rather than increase its counted work. The audit changes neither policy execution nor the frozen selections.

A.4 Primary expenditure and runtime

The primary selected configuration $(2, 2)$ has a higher success count than $(2, 5)$ on seven tasks and ties on one (Table 8). There are 24 paired wins and seven losses. In a descriptive leave-one-task-out sensitivity analysis, the equally weighted difference over the remaining seven tasks stays positive, ranging from 9.3 to 12.1 percentage points. This post-outcome analysis does not extend inference to a new task population. The supplementary paired records retain all state and seed identities needed to reproduce the primary interval.

Table 6: Descriptive primary expenditure differences, selected minus comparator at $h = 2$, including all attempts. Intervals use the paired state bootstrap conditional on the eight tasks and were added after outcomes were observed. Timing intervals summarize sampling variation in the observed records; they do not account for unmonitored runtime interference during priority completion.

Mean per attempt	$d = 2$	$d = 5$	Difference	95% interval
Reference TFLOPs	41.1	48.7	-7.6	$[-10.6, -4.5]$
Observed rollout (s)	19.0	34.3	-15.2	$[-17.2, -13.3]$
Control steps	149.0	166.3	-17.3	$[-27.9, -6.5]$

The manifests do not retain a parameter-dtype inventory or a separate vision attention-backend inventory. The attention details in Table 7 follow the pinned language/action implementation; they do not identify every runtime kernel.

A.5 Complete confirmation outcomes

The stale-input intervention is reported separately from the confirmation comparisons. It is a 40-attempt mechanism diagnostic on the discovery identities, not an additional confirmation configuration. The held-input and matched fresh-input records share all 40 initial state and observation hashes; the stale run's remote artifacts are retained with their own manifest and schedule hashes.

The four amended comparisons are reported in their documented order, using 20,000 paired state-bootstrap resamples and analysis seed 20260913. Table 1 gives their all-attempt resource means. In particular, the $(2, 2)$ minus $(5, 2)$ expenditure difference is $+25.0$ TFLOPs per attempt, with descriptive interval $[+22.8, +27.5]$. This interval describes generated-chunk arithmetic, including the different trajectory lengths; it is not an elapsed-time comparison or evidence of equivalent quality.

Table 7: Runtime and protocol shared by all four confirmation execution phases. Module versions are the loaded versions reported at evaluation time. The supplement retains source digests and distinguishes conflicting installed distribution metadata from these loaded versions.

Setting	Value
Hardware	RTX 4060 Ti, 16 GB; driver 595.84
Runtime	Python 3.12.3; PyTorch 2.7.1+cu126; CUDA 12.6; NumPy 2.2.6; MuJoCo 3.8.1
Evaluation	Batch 1; synchronous simulation; nominal 20 Hz (50 ms control period)
Episode horizons	Spatial and Object: 280 control steps; Goal: 300
Precision controls	Autocast and TF32 disabled; deterministic algorithms enabled
Attention and reuse	Eager language/action attention with float32 scores; within-generation KV cache enabled
Compilation/warm-up	Model compilation disabled; 20 warm-up generations per condition
Evaluator source	`run_paired.py`; SHA-256 prefix `516375133d030ff4` (full digest in supplement)

Table 8: Confirmation successes by task and configuration. Each task row has 20 attempts: ten paired initial states, each with two policy seeds. The total aggregates 160 attempts per configuration. The primary $(2, 2)$ setting exceeds $(2, 5)$ on seven tasks and ties on one.

	$h = 1$			$h = 2$ (primary pair)		$h = 5$	
Task	$d = 2$	$d = 5$	$d = 10$	$d = 2$	$d = 5$	$d = 2$	$d = 5$
Goal 5	14	15	14	16	14	18	16
Goal 6	18	17	17	18	15	19	18
Goal 7	20	20	20	19	18	18	18
Goal 8	19	18	18	18	14	18	15
Object 5	20	19	19	20	20	20	20
Object 6	19	18	19	20	17	18	18
Object 7	9	9	10	10	7	7	8
Object 8	18	20	18	18	17	20	20
Total (/160)	137	136	135	139	122	138	133

Table 9: Amended descriptive comparisons using all 160 attempts per setting. These contrasts were documented after collection began and are not originally prespecified or fully outcome-blind. Intervals are task-conditional, unadjusted, and do not test a quality-preservation margin.

Setting A	Setting B	A−B (pp)	95% interval (pp)
$(2, 2)$	$(1, 10)$	$+2.5$	$[-3.8, +8.1]$
$(2, 2)$	$(5, 2)$	$+0.6$	$[-6.3, +6.9]$
$(1, 2)$	$(2, 2)$	-1.3	$[-7.5, +5.0]$
$(1, 2)$	$(1, 10)$	$+1.3$	$[-5.0, +7.5]$

Measurement verification. Episode and action records capture evaluation identities, state and observation hashes, generation and denoising counts, dispatched actions, and outcomes. Independent reconciliation checks schedules, model-boundary counts, action hashes, and condition pairing across 87,364 dispatched discovery actions. The complete confirmation audit covers all 1,120 unique attempts and 171,805 dispatched actions, including exact original-plan coverage and all four execution phases. It reproduces the unchanged primary subset of 320 attempts and 50,451 actions. The raw records include two additional completed repeats with identical outcomes and action streams; these are checked but excluded from the canonical sample. A trailing interrupted attempt contains ten actions and no completed outcome, so it is recorded separately and excluded. Checkpoint, reset, noise, and environment controls further support evaluator validity; shared upstream errors remain outside this internal-consistency boundary.

Analysis provenance. The three budget-specific comparisons, identities, and primary analysis were fixed before transfer outcomes. An outcome-blind, deadline-driven scheduling amendment changed execution order; policy execution settings remained fixed and all 320 primary attempts were included. Subsequent continuation completes the remaining original identities, not an independent replication. Deduplication retains the earliest completed record under the documented phase precedence.

The original comparisons use analysis seed 20260912 in quarter-, half-, then full-native order; a separate replay resets that seed for the submitted primary interval and reproduces its endpoints. The amended contrasts use seed 20260913. Their amendment was written after collection began, when individual outcomes were visible, and is not fully outcome-blind. The simple, empirical, and compute-only baselines were likewise documented after collection began and are interpreted descriptively. Task disaggregation, leave-one-task-out sensitivity, and expenditure intervals were added after primary outcomes were observed; they do not alter the original comparison or stopping condition. No response model is refitted on confirmation outcomes.

Analysis artifact. The accompanying `primary_transfer_supplement.zip` contains the primary episode and action records, original planned identities, paired state values, calibrations, frozen predictions, runtime extract, and analysis code. Running `python primary_reporting.py -check .` reproduces the primary counts, interval, task table, and descriptive expenditure analysis without inference. Its README specifies the selector features, penalized objective, candidate set, and normalization. Machine-specific paths are omitted from the runtime extract; checkpoint weights and simulator assets are not bundled. The continuing repository additionally retains a compressed completed-confirmation snapshot and `analyze_confirmation.py`. Its `-check` mode reconciles the full journals, original schedule, runtime manifests, repeated identities, and paired blocks, then reproduces the seven-setting resource table, all original and amended comparisons, and task-level counts. Generated evidence records source, input, and snapshot hashes. The submitted supplement remains separate from these post-submission additions.

B Frozen Selection and Diagnostic Analyses

We compare two logistic response-model families. A compute-only model uses a quadratic in log normalized rate. An allocation model uses $\log h$, $\log(d/10)$, both squared terms, and their interaction. Each family is fitted with ridge penalty 1 or 10. Leave-one-task-out log loss selects the model, which is then refitted on the discovery set. No individual seed repeat is held out while another repeat from the same task remains in training.

The candidate set is $h \in \{1, 2, 3, 5\}$ crossed with $d \in \{2, 3, 5, 8, 10, 20\}$: the 12 discovery configurations and 12 untested configurations. Their calibrated costs determine eligibility; the fitted model selects the highest predicted success under each ceiling. Ties favor lower compute, then lexicographic order.

The simple rule minimizes denoising count, then chooses the shortest feasible prefix. It reproduces every learned recommendation (Table 10); transfer outcomes cannot distinguish these selectors at the tested budgets. The compute-only choices use the original ridge-10 fit. Its quarter- and half-budget choices are absent from the transfer plan, so this evaluation cannot compare those decisions on transfer tasks.

Table 10: Discovery-based selector choices. The learned choices and comparisons were frozen before transfer; additional selector baselines were documented after collection began. The empirical selector uses the measured grid; other selectors use the declared candidate set. The simple rule exactly matches the learned choices. Selected rate/native is the learned recommendation's arithmetic rate. Frozen comparators and predicted gaps appear in Table 11; the half-native row is primary.

	Selector choice (h, d)				Learned
Ceiling/native	Learned	Simple	Empirical	Compute-only	rate/native
0.25	$(5, 2)$	$(5, 2)$	$(5, 2)$	$(5, 20)$	0.171
0.50 (primary)	$(2, 2)$	$(2, 2)$	$(2, 2)$	$(2, 10)$	0.429
1.00	$(1, 2)$	$(1, 2)$	$(1, 5)$	$(1, 10)$	0.857

Table 11: All original frozen denoising tests: selected minus comparator, holding h fixed within each row. Predicted and observed gaps and intervals are percentage points. The half-native row is primary; secondary intervals are descriptive, task-conditional, and unadjusted.

Ceiling	Setting (h, d)		Success difference (pp)		
/ native	Selected	Comparator	Predicted	Observed	95% interval
0.25	$(5, 2)$	$(5, 5)$	$+6.5$	$+3.1$	$[-2.5, +8.8]$
0.50 (primary)	$(2, 2)$	$(2, 5)$	$+4.0$	$+10.6$	$[+4.4, +16.9]$
1.00	$(1, 2)$	$(1, 5)$	$+2.2$	$+0.6$	$[-5.0, +6.3]$

The selected allocation model has log loss 0.7048, compared with 0.7067 for the best compute-only model. This small difference provides no evidence of an added predictive advantage. Under the same mean natural-log binary cross-entropy, constant $p = 0.5$ scores 0.6931 and an intercept-only model fitted separately on each training fold scores 0.7170. The response-model probabilities should therefore not be read as calibrated task-transfer forecasts. The primary test evaluates a recommendation shared with simpler selectors, so it does not establish an additional benefit from model-based ranking.

AI Use Statement

We used generative AI tools to help develop the conceptual framing and hypotheses, discuss experimental design, interpret results, search literature and sources, draft and revise the manuscript, implement and review evaluation and analysis code, and prepare tables and figures. The repository records automated journal reconciliation, reproduction of the confirmation statistics, and tests comparing plotted values with retained reports. These reproducibility checks are not substitutes for author review. The authors remain responsible for the final text, research claims, code, and artifacts.